

# Leveraging locust.io for software systems performance analysis

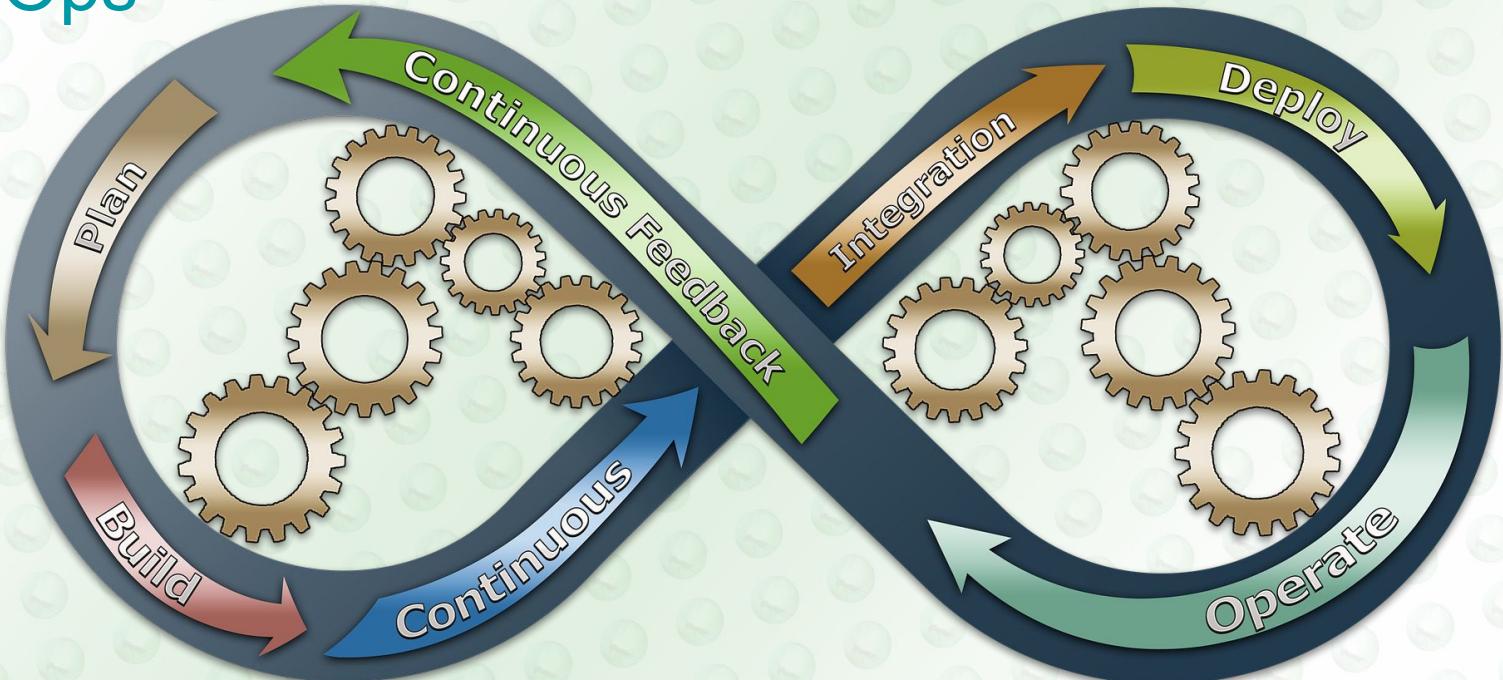
sisylabs automation smarts

# Performance Quality Operations

Culture. Infrastructure. Tools.

# Agile & DevOps

Individuals and interactions over processes and tools  
Working software over comprehensive documentation  
Customer collaboration over contract negotiation  
Responding to change over following a plan



# ...in performance engineering

- Define service quality metrics including performance based on Business requirements
- Capacity planning
- Architecture with horizontal scalability in mind
- Continuous performance testing
- Automated stress, load testing
- Static code analysis and vulnerabilities management
- Soak testing
- Dynamic scaling
- Monitoring for services and infrastructure
- Monitoring of the end user experience
- Root cause diagnostics of service disruptions

**Availability**  
**Reliability**  
**Performance**

**Accessibility**  
**Integrity**  
**Regulatory**  
**Security**

Direct and indirect impact of performance  
characteristics on Service Quality

# Cloud Systems Performance Short Version

**Throughput** - How much RPS of specific workload? :)

**Latency** - Will the response time fit in XX milliseconds?

**Resources utilization** - Is it ok to spawn a VM per X clients?

**Scalability** - X2 Throughput with X2 capacity?

**Availability** - Will users notice on X RPS when N nodes go down?

# And... Performance testing...

## Performance regression

Implement continuous performance testing for system and components

## Load

Determine the interpolated load that your production the system can handle

## Stress

Verify components durability and recovery time

# Infrastructure

- **Energy efficient computing and storage**
- High availability
- Agile workloads scheduling
- Automated operations
- Maintenance costs

...

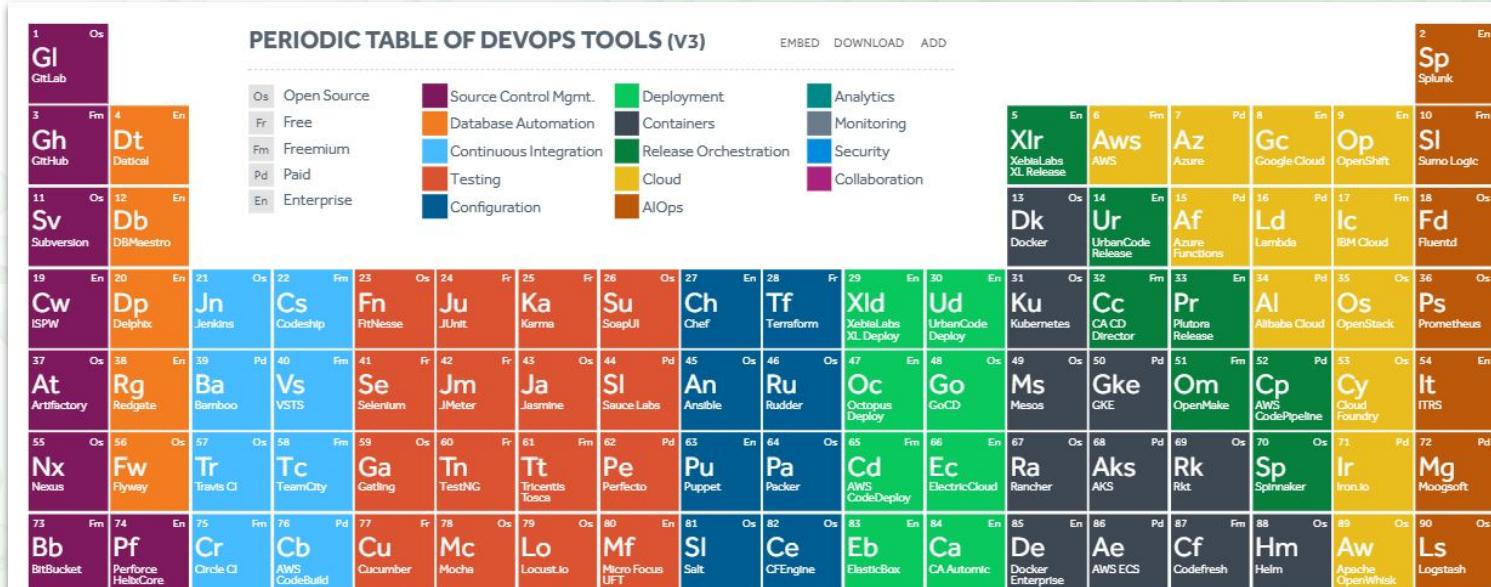


*And does not really matter if you are in cloud or on-prem.  
What really matters is **efficiency**.*

# Toolset

**PERIODIC TABLE OF DEVOPS TOOLS (V3)**

EMBED DOWNLOAD ADD



The Periodic Table of DevOps Tools (V3) is a comprehensive guide to the landscape of DevOps tools. It is organized into a grid where each cell contains a tool name and its corresponding details. The columns represent atomic numbers (1 to 120), and the rows represent periodic groups. The tools are categorized into several groups based on their primary function:

- Source Control Mgmt.**: GitLab, GitHub, Databatic, Subversion, DBMaestro, ISPW, Delphix, Jenkins, CodeShip, RnTesse, JUnit, Karma, SoapUI, Chef, Terraform, XLab XL Release, AWS, Azure, GoogleCloud, OpenShift, Splunk, Sumo Logic.
- Deployment**: Docker, UrbanCode Release, Kubernetes, CA CD Director, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Analytics**: XLab XL Release, AWS, Azure, GoogleCloud, OpenShift, Splunk, Sumo Logic.
- Database Automation**: Octopus Deploy, Sauce Labs, Ansible, RuRudder, Puppet, Packer, AWS CodeDeploy, ElectricCloud, Rancher, AKS, Rkt, Spinaker, Iron.io, Moosoft.
- Containers**: Docker, UrbanCode Release, Kubernetes, CA CD Director, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Monitoring**: UrbanCode Release, Kubernetes, CA CD Director, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Continuous Integration**: Jenkins, Bamboo, VSTS, Selenium, JMeter, Jasmine, Sauce Labs, Ansible, RuRudder, Octopus Deploy, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Testing**: Bamboo, VSTS, Selenium, JMeter, Jasmine, Sauce Labs, Ansible, RuRudder, Octopus Deploy, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Release Orchestration**: Puppet, Packer, AWS CodeDeploy, ElectricCloud, Rancher, AKS, Rkt, Spinaker, Iron.io, Moosoft.
- Cloud**: XLab XL Release, AWS, Azure, GoogleCloud, OpenShift, Splunk, Sumo Logic.
- Security**: Docker, UrbanCode Release, Kubernetes, CA CD Director, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Collaboration**: Jenkins, Bamboo, VSTS, Selenium, JMeter, Jasmine, Sauce Labs, Ansible, RuRudder, Octopus Deploy, GoCD, Mesos, GKE, OpenMake, AWS CodePipeline, Cloud Foundry, Alibaba Cloud, OpenStack, Prometheus.
- Configuration**: CircleCI, AWS CodeBuild, Cucumber, Mocha, Locust.io, Micro Focus UFT, Salt, CFEngine, ElasticBox, CA Automic, Docker Enterprise, AWS ECS, Codefresh, Helm, Apache OpenWhisk, Logstash.

**Legend:**

- Os**: Open Source
- Fr**: Free
- Fm**: Freemium
- Pd**: Paid
- En**: Enterprise

**Tool Details:**

- Xli**: XebiaLabs XL Impact
- Ki**: Kibana
- Nr**: New Relic
- Dt**: Dynatrace
- Dd**: Datadog
- Ad**: AppDynamics
- El**: Elasticsearch
- Ni**: Nagios
- Zb**: Zabbix
- Zn**: Zenoss
- Cx**: Checkmark SAST
- Sg**: Signal Sciences
- Bd**: BlackDuck
- Sr**: SonarQube
- Hv**: HashiCorp Vault
- Sw**: ServiceNow
- Jr**: Jira
- Tl**: Trello
- Sl**: Slack
- St**: Stride
- Cn**: CollabNet VersionOne
- Ry**: Remedy
- Ac**: Agile Central
- Og**: OpsGuru
- Pd**: Pagerduty
- Sn**: Snort
- Tw**: Tripwire
- Ck**: CyberArk Conjur
- Vc**: Veracode
- Ff**: Fortify SCA

**Navigation:**

- Follow @xebialabs**
- Publication Guidelines**
- Download**

# Locust.io

# Web UI

The screenshot shows the Locust web interface running on a host under load. The top bar displays the Locust logo, the host URL (http://hostunderload), and the status as "RUNNING 2000 users". It also shows 6 slaves, an RPS of 10861.5, and 0% failures. There are buttons for "STOP" and "Reset Stats". Below the header, a navigation bar includes links for Statistics, Charts, Failures, Exceptions, Download Data, and Slaves. The main content area is titled "Statistics" and contains a table with the following data:

Type	Name	# requests	# fails	Median (ms)	Average (ms)	Min (ms)	Max (ms)	Content Size (bytes)	# reqs/sec
GET	Get asset	116410	0	120	112	69	2781	612	3618.9
GET	List projects	116854	0	86	136	0	3034	612	3618.1
GET	Retrieve project information	115983	0	110	108	95	2550	612	3624.5
<b>Total</b>		<b>349247</b>	<b>0</b>	<b>110</b>	<b>118</b>	<b>0</b>	<b>3034</b>	<b>612</b>	<b>10861.5</b>

# Command line execution

Name	# reqs	# fails	Avg	Min	Max		Median	req/s
GET Get asset	176466	0(0.00%)	1	0	20		1	4356.30
GET List projects	176476	0(0.00%)	1	0	27		1	4356.70
GET Retrieve project information	176454	0(0.00%)	1	0	19		1	4355.50
Total	529396	0(0.00%)					13068.50	
Name	# reqs	# fails	Avg	Min	Max		Median	req/s
GET Get asset	189343	0(0.00%)	1	0	20		1	4349.00
GET List projects	189355	0(0.00%)	1	0	27		1	4349.40
GET Retrieve project information	189331	0(0.00%)	1	0	19		1	4348.50
Total	568029	0(0.00%)					13046.90	
Name	# reqs	# fails	Avg	Min	Max		Median	req/s
GET Get asset	200988	0(0.00%)	1	0	21		1	4310.00
GET List projects	201003	0(0.00%)	1	0	27		1	4310.60
GET Retrieve project information	200979	0(0.00%)	1	0	19		1	4309.90
Total	602970	0(0.00%)					12930.50	
Name	# reqs	# fails	Avg	Min	Max		Median	req/s
GET Get asset	200988	0(0.00%)	1	0	21		1	4310.00
GET List projects	201003	0(0.00%)	1	0	27		1	4310.60
GET Retrieve project information	200979	0(0.00%)	1	0	19		1	4309.90
Total	602970	0(0.00%)					12930.50	
Name	# reqs	# fails	Avg	Min	Max		Median	req/s
GET Get asset	212632	0(0.00%)	1	0	21		1	4170.40
GET List projects	212645	0(0.00%)	1	0	27		1	4170.70
GET Retrieve project information	212613	0(0.00%)	1	0	19		1	4170.20
Total	637890	0(0.00%)					12511.30	

# Distributed

## **Web UI:**

```
locust -f locustfile.py --slave --master-host=0.0.0.0 &> locust_slave.log
```

```
locust -f locustfile.py --master --host=http://hostunderload
```

## **Console:**

```
locust -f locustfile.py --slave --master-host=0.0.0.0 &> locust_slave.log
```

```
locust -f locustfile.py -c 5000 --master --host=http://hostunderload --expect-slaves  
2 --no-web
```

# HttpLocust

```
locustfile_user.py x

1  from locust import Locust, HttpLocust, TaskSet, task
2
3  # TaskSet is a collection of tasks a user (Locust) should
4  class DemoHttpTaskSet(TaskSet):
5      @task
6      def index(self):
7          self.client.get("/") # session-aware client provided by HttpLocust
8
9  class DemoTaskSet(TaskSet):
10     @task
11     def Demo_task(self):
12         print("Locust instance (%r) does a specific task" % (self.locust))
13
14 class DemoHttpLocust(HttpLocust): # HttpLocust - user with HTTP client (self.client)
15     weight = 2 # DemoHttpLocust will be spawned twice more often than DemoLocust
16     task_set = DemoHttpTaskSet
17     # minimum and maximum time in milliseconds, that a simulated user will wait
18     # between executing each task (randomized)
19     min_wait = 1000
20     max_wait = 3000
21
22     def setup(self):
23         print("(%r) Locust setup - first action")
24
25     def teardown(self):
26         print("(%r) Locust teardown - last action")
27
28 class DemoLocust(Locust): # Locust class represents 1 simulated user
29     weight = 1 # DemoLocust will be spawned twice less often than DemoHttpLocust
30     # Tasks that user needs to execute
31     task_set = DemoTaskSet
```

# TaskSet

```
locustfile_taskset.py ×

 1  from locust import HttpLocust, TaskSet, task
 2  import datetime
 3  class DemoHttpTaskSet(TaskSet):
 4      def setup(self):
 5          print("%r) DemoHttpTaskSet setup (once per class)" % datetime.datetime.now())
 6
 7      def teardown(self):
 8          print("%r) DemoHttpTaskSet teardown (once per class)" % datetime.datetime.now())
 9
10     def on_start(self):
11         print("%r) DemoHttpTaskSet on_start" % datetime.datetime.now())
12
13     def on_stop(self):
14         print("%r) DemoHttpTaskSet on_stop" % datetime.datetime.now())
15
16     @task(3)
17     def demo_task(self):
18         self.client.get("/", name="Demo Task")
19
20     @task(1)
21     class DemoHttpSubTaskSet(TaskSet):
22         # setup teardown and on_start also applicable
23         @task
24         def demo_sub_task(self):
25             self.client.get("/", name="Demo Sub Task")
26
27         @task
28         def stop(self):
29             self.interrupt() # Required to terminate the TaskSet execution
30
31     class DemoHttpLocust(HttpLocust):
32         task_set = DemoHttpTaskSet
```

# TaskSequence

```
locustfile_sequence.py x

1  from locust import HttpLocust, TaskSequence, seq_task, task
2
3  class DemoTaskSequence(TaskSequence):
4      # seq_task - Defines the order of task execution
5      # login -> get_info -> get_asset
6      @seq_task(1)
7      def login(self):
8          self.client.get("/", name="Login")
9
10     @seq_task(2)
11     def get_info(self):
12         self.client.get("/", name="Get info")
13
14     @seq_task(3)
15     @task(2) # Defines that the get asset runs twice
16     def get_asset(self):
17         self.client.get("/", name="Get asset")
18
19     class DemoHttpLocust(HttpLocust):
20         task_set = DemoTaskSequence
```

# IDE debugging

```
1  from locust import HttpLocust, TaskSet, task
2
3  class DemoHttpTaskSet(TaskSet):
4      @task
5      def demo_task(self):
6          res = self.client.get("/", name="Debugging Example")
7          print(str(res))
8
9  class DemoHttpLocust(HttpLocust):
10     task_set = DemoHttpTaskSet
11
12 if __name__ == '__main__':
13     from argparse import Namespace
14     from locust import runners as r
15     opts = Namespace()
16     opts.host = "http://google.com"
17     opts.num_clients = 1
18     opts.hatch_rate = opts.num_clients
19     opts.reset_stats=False
20     r.locust_runner = \
21         r.LocalLocustRunner(locust_classes=[DemoHttpLocust],
22                             options=opts)
23     r.locust_runner.start_hatching(wait=True)
24     r.locust_runner.greenlet.join()
```

# HTML Parsing

```
locustfile_html.py x

1  from locust import HttpLocust, TaskSequence, seq_task, task
2  from pyquery import PyQuery #pip install pyquery
3  import random
4
5  class DemoTaskSequence(TaskSequence):
6      @seq_task(1)
7      @task(1)
8      def index_page(self):
9          r = self.client.get("/", name="Get Root")
10         pq = PyQuery(r.content)
11         self.links = pq("a") # Get links
12
13     @seq_task(2)
14     @task(1)
15     def load_page(self, url=None):
16         url = random.choice(self.links).attrib["href"]
17         print(str(url))
18         self.client.get(url)
19
20     class DemoHttpLocust(HttpLocust):
21         task_set = DemoTaskSequence
22         host = "https://github.com/"
```

# Hacking Options

- **Custom Clients**
- Performance tuning with **FastHTTPLocust**
- **Extending the API**
- Anything... It's just **Python**

# Typical implementation

If you need more than just a locustfile.py

# Solution structure

```
locust-training
  common
    __init__.py
    config.py
    logger.py
  docs
  fixtures
    __init__.py
    shop.py
  locustfiles
    shop.py
  .dockerignore
  .gitignore
  Dockerfile
  README.md
```

- Common utility libraries (format converters, custom clients)
- Fixtures reused in multiple tests to generate test data
- Locustfiles
- Docker and git configuration

# Building multi-scenario execution

 Codename **Taurus**

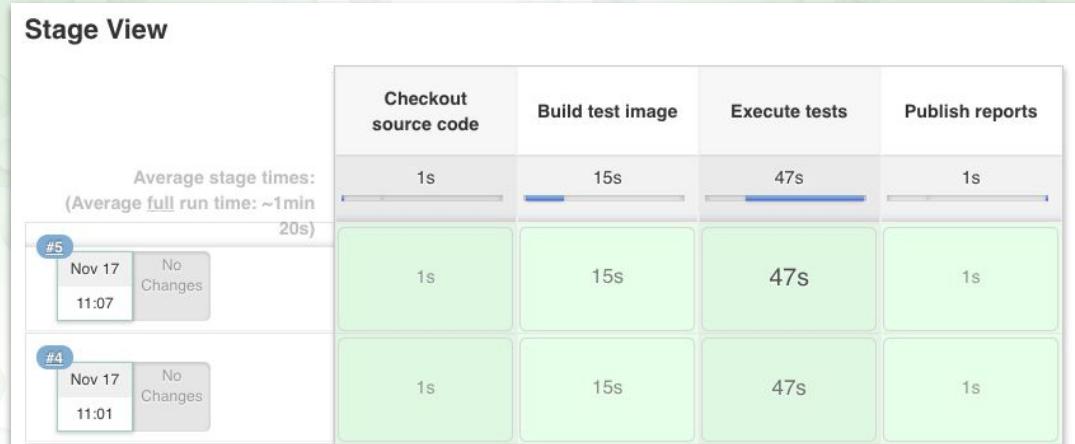
```
! hello.yaml x
1  execution:
2    - executor: locust
3      concurrency: 2
4      ramp-up: 10s
5      iterations: 10
6      scenario: hello
7      slaves: 2
8    - executor: locust
9      concurrency: 5
10     ramp-up: 5s
11     iterations: 10
12     scenario: hello
13     slaves: 4
14  ###
15  scenarios:
16    hello:
17      default-address: http://google.com
18      script: hello.py
```



pip install bzt

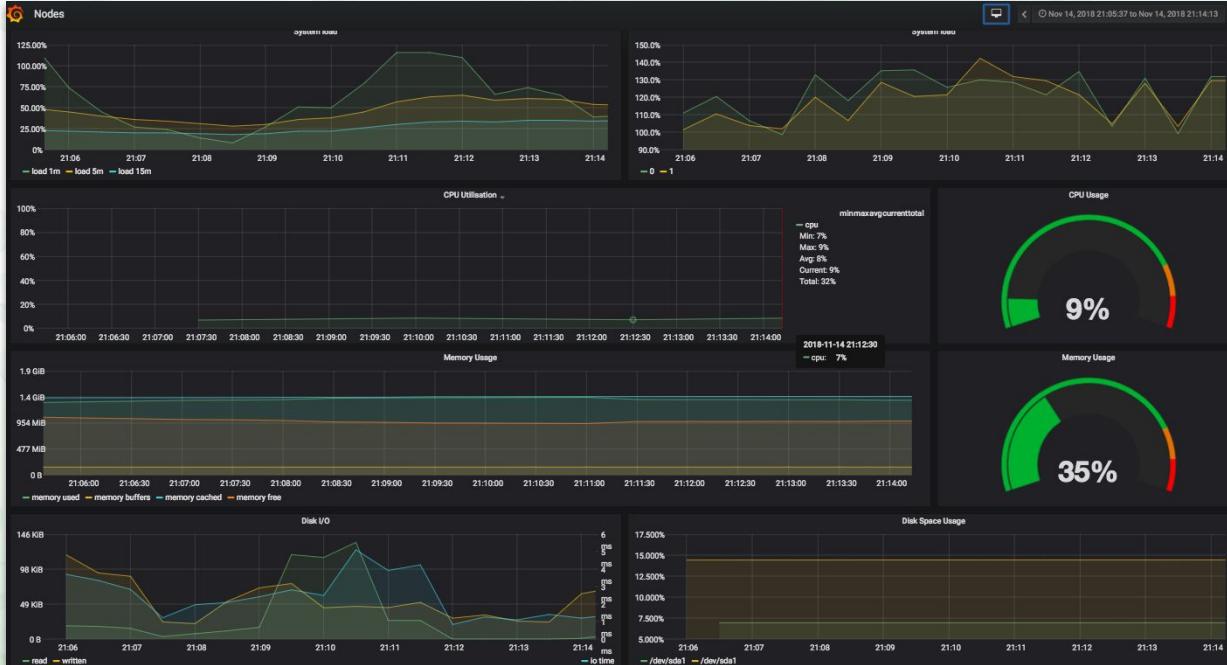
# Continuous performance testing

- Use a continuous integration system to run the performance tests
- Track the performance results within CI or with external services



# Performance issues analysis toolset

# Monitoring



 Prometheus

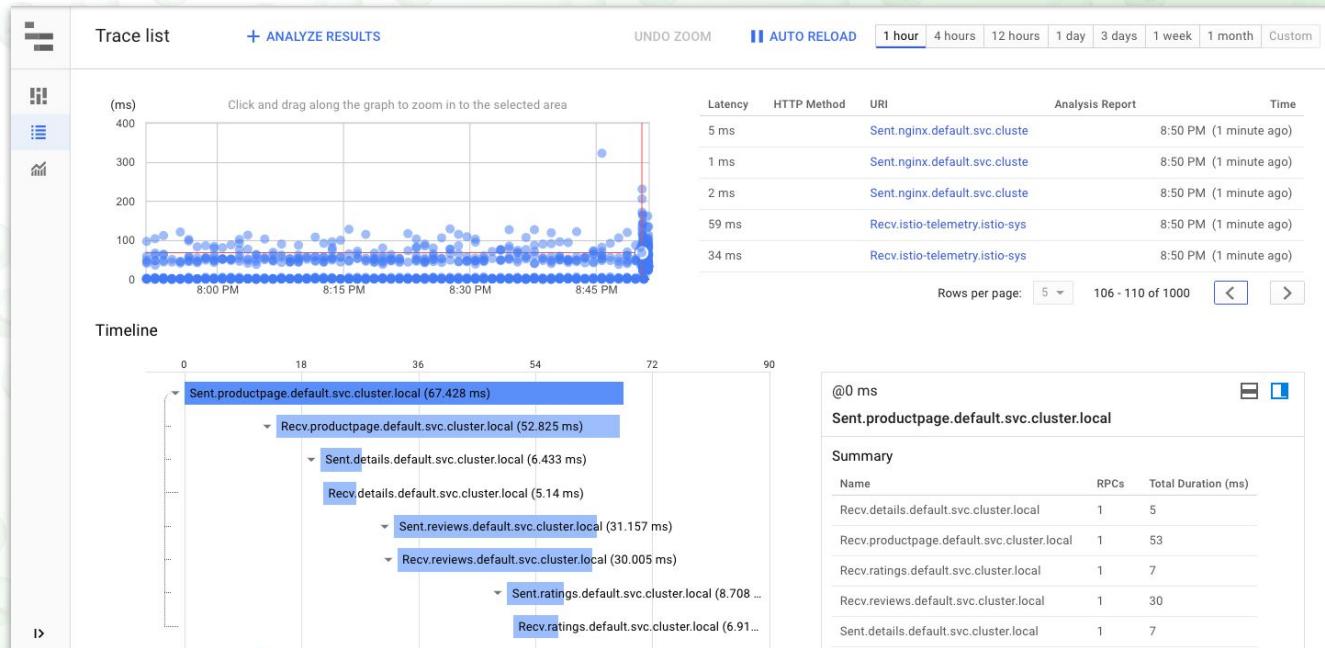
 influxdb

 Nagios®

 sensu Stackdriver

 DATADOG

# Distributed tracing



# Applications performance monitoring

OpenCensus

Insping

New Relic APM

AppDynamics

Foglight

ExtraHop



\* Image source: <https://newrelic.com/>

# Centralized logging



**kibana**

169,448 hits      New    Save    Open    Share    Reporting    Auto-refresh    Last 15 minutes

Search... (e.g. status:200 AND extension:PHP)      Options   

**Discover**

**Visualize**

**Dashboard**

**Timeline**

**Machine Learning**

**APM**

**Dev Tools**

**Monitoring**

**Management**

**Guest User**

Add a filter

Selected Fields: filebeat-\*

Available Fields:

Count

November 16th 2018, 22:58:26.454 - November 16th 2018, 23:13:26.454 — Auto

Time @timestamp per 30 seconds

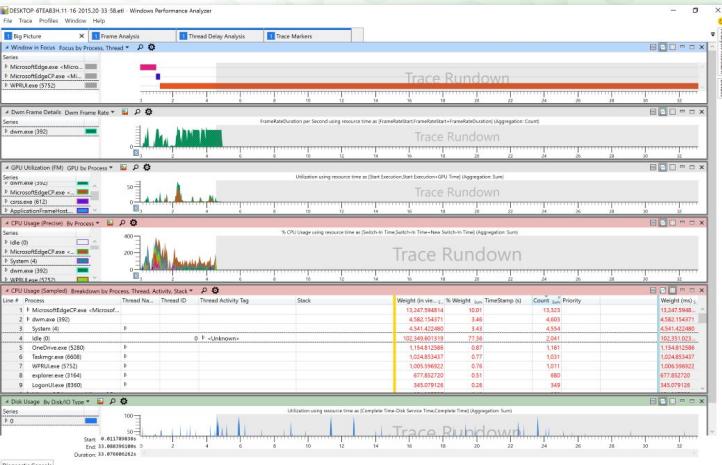
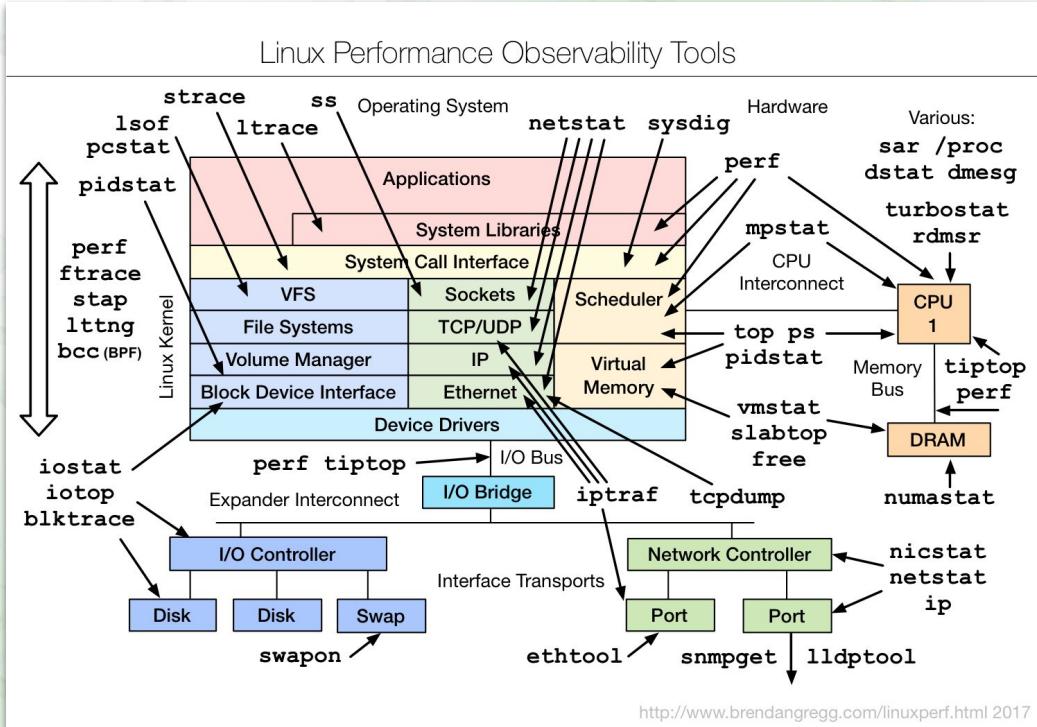
Time	Count
22:59:00	~6,000
23:01:00	~6,500
23:03:00	~5,800
23:05:00	~6,200
23:07:00	~7,000
23:09:00	~6,000
23:11:00	~5,500

Time \_source

▶ November 16th 2018, 23:13:24.000 kubernetes.container.name: nginx-ingress-controller  
mo-elastic--default-pool-1-8c155651-v29q kubernetes.namespace: default kub  
2950176807 kubernetes.labels.k8s-app: nginx-ingress-  
nginx.access.referrer: https://demo.elastic.co/app/m

▶ November 16th 2018, 23:13:24.000 kubernetes.container.name: nginx-ingress-controller  
mo-elastic--default-pool-1-8c155651-v29q kubernetes.namespace: default kub  
2950176807 kubernetes.labels.k8s-app: nginx-ingress-  
nginx.access.referrer: https://demo.elastic.co/app/m

# OS Performance analysis



# Summary

What's next?

# Performance engineer

Performance engineer job is not to write a script and break the system with many RPS, but it's to control the software delivery process culture, ensure that core quality gates are in place and are effective.

Review the toolset that is essential in performance engineering:

- Monitoring/APM
- Distributed tracing
- Logging
- Analytics

# Locust ...

- **IS** a good tool to emulate complex user workflows using plain Python scenarios
- **IS** horizontally scalable and allows simulating 50k+ RPS in distributed mode
- **IS** customizable
- **IS NOT** the tool for generating max density RPS for simple GET requests

# Q&A

# References

Name	Reference link
The Art of Application Performance Testing, 2nd Edition	<a href="http://shop.oreilly.com/product/0636920033233.do">http://shop.oreilly.com/product/0636920033233.do</a>
<a href="https://github.com/GoogleCloudPlatform/distributed-load-testing-using-kubernetes">https://github.com/GoogleCloudPlatform/distributed-load-testing-using-kubernetes</a>	<a href="https://github.com/GoogleCloudPlatform/distributed-load-testing-using-kubernetes">https://github.com/GoogleCloudPlatform/distributed-load-testing-using-kubernetes</a>
Linux performance observability	<a href="https://medium.com/@chrishantha/linux-performance-observability-tools-19ae2328f87f">https://medium.com/@chrishantha/linux-performance-observability-tools-19ae2328f87f</a>
Windows performance analyzer	<a href="https://docs.microsoft.com/en-us/windows-hardware/test/wpt/windows-performance-analyzer">https://docs.microsoft.com/en-us/windows-hardware/test/wpt/windows-performance-analyzer</a>

Name	Reference link
Examples for this training	<a href="https://github.com/runonautomation/locust-training"><u>https://github.com/runonautomation/locust-training</u></a>
Example solution structure	<a href="https://github.com/runonautomation/locust-framework"><u>https://github.com/runonautomation/locust-framework</u></a>
Locust documentation	<a href="https://docs.locust.io/en/stable/"><u>https://docs.locust.io/en/stable/</u></a>
Locust custom clients	<a href="https://medium.com/locust-io-lets-get-some-fun/locust-custom-client-23e205f4611f"><u>https://medium.com/locust-io-lets-get-some-fun/locust-custom-client-23e205f4611f</u></a>
Extending locust monitoring	<a href="https://www.blazemeter.com/blog/locust-monitoring-with-grafana-in-just-fifteen-minutes"><u>https://www.blazemeter.com/blog/locust-monitoring-with-grafana-in-just-fifteen-minutes</u></a>
Locust experiments	<a href="https://medium.com/locust-io-experiments"><u>https://medium.com/locust-io-experiments</u></a>